


Latency: repo methodology vs true kernel time



Useful FLOP throughput: max ~4.6 TFLOPS = 0.5% of peak (989 TF)
neither kernel is compute-throughput bound; INT4's high SM% is dequant overhead

