


W8A8 (INT8 tensor-core MMA, scales on accumulator only) vs cuBLAS FP16 vs Triton W4A16
MLA reconstruction BMM, H100, CUDA-graph timed; W8A8 includes dynamic activation quant

A — bs=1 decode across the L2 cliff:
first quantized kernel that beats cuBLAS FP16



B — at the 16 MB L2-resident MLA shape no quantized kernel wins at any bs (L2 barrier confirmed causally)

