


A — tuned-CUDA MoE latency (Mixtral E=8,H=4096,l=14336,topk=2)



B — CUDA quant crosses below 1.0 at ~600 tok; Triton did not

