


A — sparse MLA decode latency (C=4096, topk=2048, h_q=128)



B — FP8 KV washes out at high batch (throughput regime)

