


The capacity gate: quantization pays only above the measured effective L2 capacity (NVIDIA H100 80GB HBM3, clocks locked)

W8A8 (INT8 MMA, no in-core dequant)



W4A16 (Triton in-core dequant)

